

Analyze A/B Test Results

This project will assure you have mastered the subjects covered in the statistics lessons. We have organized the current notebook into the following sections:

- Introduction
- Part I - Probability
- Part II - A/B Test
- Part III - Regression
- Final Check
- Submission

Specific programming tasks are marked with a **ToDo** tag.

Introduction

A/B tests are very commonly performed by data analysts and data scientists. For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should:

- Implement the new webpage.
- Keep the old webpage, or
- Perhaps run the experiment longer to make their decision.

Each **ToDo** task below has an associated quiz present in the classroom. Though the classroom quizzes are **not necessary** to complete the project, they help ensure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the **final** student checklist.

Tip: Though not a **hardcore** statistics, students can attempt the classroom quizzes to ensure statistical numeric values are calculated correctly in many cases.

Part I - Probability

To get started, let's import our libraries.

```
In [15]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
# We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

ToDo 1.1

Now, read in the `ab_data.csv` data. Store it in `df`. Below is the description of the data, there are a total of 5 columns:

Data column	Purpose	Valid values
user_id	Unique ID	refers values
timestamp	Time stamp when the user visited the webpage	-
group	In the current A/B experiment, the users are categorized into two total groups. The <code>control</code> group users are expected to be served with <code>old_page</code> , and <code>treatment</code> group users are matched with the <code>new_page</code> . However, some inaccurate rows are present in the real data, such as a <code>control</code> group user is matched with a <code>new_page</code> .	<code>['control', 'treatment']</code>
landing_page	It denotes whether the user visited the old or new webpage	<code>['old_page', 'new_page']</code>
converted	It denotes whether the user decided to pay for the company's product. Here, <code>1</code> means yes, the user bought the product.	<code>[0, 1]</code>

<center>Use your dataframe to answer the questions in Quiz 1 of the classroom.

a. Read in the dataset from the `ab_data.csv`. It's and take a look at the top few rows here:

```
In [16]: df=pd.read_csv('ab_data.csv')
```

b. Use the cell below to find the number of rows in the dataset.

```
In [17]: print(df.shape)
(284878, 5)
```

The number of unique users in the dataset.

```
In [18]: df.user_id.unique()
print(df.user_id.nunique())
(290584, )
```

d. The proportion of users converted.

```
In [19]: df.query('converted==1').user_id.nunique()/df.query('converted').count()
```

```
Out[19]: user_id      0.0
         timestamp  0.0
         landing_page 0.0
         converted   0.0
         dtype: float64
```

e. The number of times when the "group" is `'treatment'` but `'landing_page'` is not a `'new_page'`.

```
In [19]: df2=df.dropna()
df2[(df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')] == False] shape[0]
```

Out[19]: 3893

f. Do any of the rows have missing values?

```
In [111]: df.isnull().sum()
```

```
Out[111]: user_id      0
         timestamp  0
         group      0
         landing_page 0
         converted  0
         dtype: int64
```

ToDo 1.2

In a particular row, the **group** and **landing_page** columns should have either of the following acceptable values:

user_id	timestamp	group	landing_page	converted
xxxx	xxxx	control	old_page	x
xxxx	xxxx	treatment	new_page	x

It means, the `control` group users should match with `old_page`, and `treatment` group users should matched with the `new_page`.

However, for rows where `treatment` does not match with `new_page` or `control` does not match with `old_page`, we cannot be sure if such rows truly received the new or old webpage.

Use **Quiz 2** in the classroom to figure out how should we handle the rows where the group and landing_page columns don't match?

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in `df2`.

```
In [112]: # Remove rows where the group and landing_page columns do not match
df2 = df.drop(df[(df.group == 'treatment') & (df.landing_page == 'old_page')].index)
df2 = df2.drop(df[(df.group == 'control') & (df.landing_page == 'new_page')].index)
```

```
Out[112]: C:\Users\kheoh\AppData\Local\Temp\ipykernel_2684\56285488.py:4: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
df2 = df2.drop(df[(df.group == 'control') & (df.landing_page == 'new_page')].index)
```

c. Double-check all of the incorrect rows were removed from `df2` -

```
In [113]: # Output of the statement below should be 0
df2[(df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')] == False] shape[0]
```

Out[113]: 0

ToDo 1.3

Use `df2` and the cells below to answer questions for **Quiz 3** in the classroom.

a. How many unique user_ids are in `df2`?

```
In [114]: df2.user_id.unique()
print(df2.user_id.nunique())
```

Out[114]: (290584,)

b. There is one user_id repeated in `df2`. What is it?

```
In [115]: df2[df2.duplicated('user_id', keep=False)]
```

```
Out[115]: user_id      timestamp      group      landing_page      converted
2892    771802    2017-01-14 02:55:59.500327    treatment      new_page          0
```

c. Display the rows for the duplicate user_id?

```
In [116]: df2[df2.duplicated('user_id', keep=False)]
```

```
Out[116]: user_id      timestamp      group      landing_page      converted
1889    771802    2017-01-09 05:37:56.761000    treatment      new_page          0
2892    771802    2017-01-14 02:55:59.500327    treatment      new_page          0
```

d. Remove one of the rows with a duplicate user_id from the `df2` dataframe.

```
In [117]: # Remove one of the rows with a duplicate user_id...
df2 = df2.drop(df2[df2.user_id == df2.user_id].index]
```

```
Out[117]: user_id      timestamp      group      landing_page      converted
0      851104    2017-01-21 22:11:48.556739    control      old_page          0
1      804228    2017-01-12 08:01:45.199739    control      old_page          0
2      804228    2017-01-12 08:01:45.199739    treatment    new_page          0
3      80584    2017-01-09 18:28:03.147656    treatment    new_page          0
4      864975    2017-01-21 01:52:26.210827    control      old_page          1
...      ...      ...      ...      ...      ...
284675    751911    2017-01-03 22:29:36.630929    control      old_page          0
284676    845632    2017-01-22 00:57:57.070732    control      old_page          0
284677    734069    2017-01-22 11:45:03.493454    control      old_page          0
284678    697124    2017-01-15 01:20:28.957438    control      old_page          0
284677    715931    2017-01-16 12:40:24.467437    treatment    new_page          0
```

290584 rows x 5 columns

ToDo 1.4

Use `df2` in the cells below to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

Tip: The probability you'll compute represents the overall "converted" success rate in the population and you may call it $p_{\text{population}}$.

a. What is the conversion rate for p_{new} under the null hypothesis?

```
In [118]: df2['converted'].mean()
```

```
Out[118]: 0.11959768724499628
```

b. Given that an individual was in the `'control'` group, what is the probability they converted?

```
In [119]: df2.groupby('group')['converted'].mean()
```

```
Out[119]: group
control    780164.072694    0.120386
treatment   797845.718260    0.118888
```

c. Given that an individual was in the `'treatment'` group, what is the probability they converted?

```
In [120]: df2.groupby('group')['converted'].mean()
```

```
Out[120]: group
control    780164.072694    0.120386
treatment   797845.718260    0.118888
```

Calculate the actual difference (`obs_diff`) between the conversion rates for the two groups. You will need that later.

```
In [121]: # Calculate the actual difference (obs_diff) between the conversion rates for the two groups.
obs_diff = 0.120386 - 0.118888
print(obs_diff)
```

Out[121]: 0.00149788889000081

d. What is the probability that an individual received the new page?

```
In [122]: len(df2.query('landing_page == "new_page"'))
new_page_probabilility = 145311/290585
```

Out[122]: 0.500036646742886

e. Consider your results from parts (a) through (d) above, and explain below whether the new `'treatment'` group users lead to more conversions.

between a and d there is a little difference as the probability of an individual converting regardless of the page they receive=0.119 and the probability that an individual received the new page=0.5

Part II - A/B Test

Since a timestamp is associated with each event, you could run a hypothesis test continuously as long as you observe the events.

However, then the hard questions would be:

- Do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time?
- How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

ToDo 2.1

For now, consider you need to make the decision just based on all the data provided.

Tip: Recall that you just calculated that the "converted" probability (or rate) for the old page is slightly higher than that of the new page (ToDo 1.4).

If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should be your null and alternative hypotheses (H_0 and H_1)?

You can state your hypothesis in terms of words or in terms of μ_{old} and μ_{new} , which are the "converted" probability (or rate) for the old and new pages respectively.

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is: do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should be your null and alternative hypotheses be?

and the converted rates for the old and new pages respectively.

2. Assume under the null hypothesis, and both have "true" success rates equal to the converted success rate regardless of page - that is and are equal. Furthermore, assume they are equal to the converted rate in `ab_data.csv` regardless of the page.

ToDo 2.2 - Null Hypothesis H_0 Testing

Under the null hypothesis H_0 , assume that μ_{new} and μ_{old} are equal. Furthermore, assume that μ_{new} and μ_{old} both are equal to the **converted** success rate in the `df2` data regardless of the page. So, our assumption is:

$$p_{new} = p_{old} = p_{\text{population}}$$

In this section, you will:

- Simulate (bootstrap) sample data set for both groups, and compute the "converted" probability p for each samples.
- Use a sample size for each group equal to the ones in the `df2` data.
- Compute the difference in the "converted" probability for the two samples above.

Perform the sampling distribution for the "difference in the converted probability" between the two simulated-samples over 10,000 iterations; and calculate an estimate.

Use the cells below to provide the necessary parts of this simulation. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the conversion rate for p_{new} under the null hypothesis?

```
In [123]: p_new = df2.converted.mean()
print(p_new)
```

```
Out[123]: 0.11959768724499628
```

b. What is the conversion rate for p_{old} under the null hypothesis?

```
In [124]: p_old = df2['converted'].mean()
print(p_old)
```

```
Out[124]: 0.11959768724499628
```

c. What is n_{new} , the number of individuals in the treatment group?

Hint: The treatment group users are shown the new page.

```
In [125]: n_new=df2.query('landing_page=="new_page"').shape[0]
print(n_new)
```

```
Out[125]: 145310
```

d. What is n_{old} , the number of individuals in the control group?

```
In [126]: n_old=df2.query('landing_page=="old_page"').shape[0]
print(n_old)
```

```
Out[126]: 145274
```

e. Simulate Sample for the Treatment Group
Simulate n_{new} transactions with a conversion rate of p_{new} under the null hypothesis.

Hint: Use `numpy.random.choice()` method to randomly generate n_{new} number of values. Store these n_{new} 1's and 0's in the `new_page_converted` numpy array.

```
In [127]: # Simulate a Sample for the Treatment Group
new_page_converted = np.random.binomial(1, p_new, n_new) # bootstrapping
old_page_converted = np.random.binomial(1, p_old, n_old) # bootstrapping
p_diffs.append(new_page_converted.mean() - old_page_converted.mean())
```

```
Out[127]: 0.11843644621842957
```

f. Simulate Sample for the Control Group
Simulate n_{old} transactions with a conversion rate of p_{old} under the null hypothesis.

Store these n_{old} 1's and 0's in the `old_page_converted` numpy array.

```
In [128]: # Simulate a Sample for the Control Group
old_page_converted = np.random.binomial(1, p_old, n_old)
old_page_converted.mean()
```

```
Out[128]: 0.1190965666061374
```

g. Find the difference in the "converted" probability ($p'_{\text{new}} - p'_{\text{old}}$) for your simulated samples from the parts (e) and (f) above.

```
In [129]: new_page_converted.mean() - old_page_converted.mean()
```

```
Out[129]: -0.0088573164421841732
```

h. Sampling distribution
Re-create `new_page_converted` and `old_page_converted` and find the ($p'_{\text{new}} - p'_{\text{old}}$) value 10,000 times using the same simulation process you used in parts (a) through (g) above.

Store all ($p'_{\text{new}} - p'_{\text{old}}$) values in a NumPy array called `p_diffs`.

```
In [130]: # sampling distribution
p_diffs = []
for i in range(10000):
    new_page_converted = np.random.binomial(1, p_new, n_new) # bootstrapping
    old_page_converted = np.random.binomial(1, p_old, n_old) # bootstrapping
    p_diffs.append(new_page_converted.mean() - old_page_converted.mean())
```

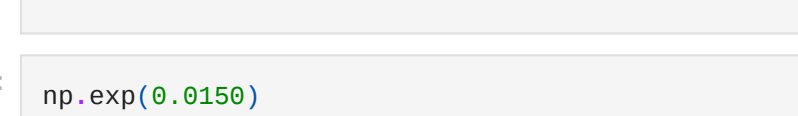
i. Histogram
Plot a histogram of the `p_diffs`. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

Also, use `plt.axvline()` method to mark the actual difference observed in the `df2` data (recall `obs_diff`). In the chart.

```
In [132]: p_diffs = plt.axvline((obs_diff))
```

```
Out[132]: # plot sampling distribution
plt.hist(p_diffs)
plt.xlabel('p_diffs')
plt.ylabel('trials')
plt.title('Difference between old page/new page after many trials')
```

```
Text(0.5, 1.5, 'Difference between old page/new page after many trials')
```



j. What proportion of the `p_diffs` are greater than the actual difference observed in the `df2` data?

```
In [132]: actual_diff = df2.converted[df2.group == 'treatment'].mean() - df2.converted[df2.group == 'control'].mean()
(actual_diff = p_diffs.mean())
```

```
Out[132]: 0.9885
```

k. Please explain in words what you have just computed in part j above.

- What is the value called in scientific context?
- What does the value signify in terms of whether or not there is a difference between the new and old pages? Hint: Compare the value above with the "Type I error rate (0.05)".

We observed difference in means of converted old page and converted new page. we had random choices of these mean converted values for the observed event. The actual difference was calculated from the dataset ab_data.csv value. In difference in means, the p-value. Our p-value is the same as the critical value of 0.05 so we fail to reject the null hypothesis. new page couldn't attract user more than old page

Using Built-in Methods for Hypothesis Testing

We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance.

Fill in the statements below to calculate the:

- `convert_old`: number of conversions with the old_page
- `convert_new`: number of conversions with the new_page
- `n_old`: number of individuals who were shown the old_page
- `n_new`: number of individuals who were shown the new_page

```
In [133]: import statsmodels.api as sm
```

```
# number of conversions with the old_page
convert_old = df2.query('group == "control" & converted == 1')['converted'].count()
```

```
# number of conversions with the new_page
convert_new = df2.query('group == "treatment" & converted == 1')['converted'].count()
```

```
# number of individuals who were shown the old_page
n_old = df2.query('landing_page=="old_page"').shape[0]
```

```
# number of individuals who received new_page
n_new = df2.query('landing_page=="new_page"').shape[0]
```

```
convert_old, convert_new, n_old, n_new
```

```
Out[133]: (27489, 27264, 145274, 145310)
```

m. Now use `sm.stats.proportions_ztest()` method arguments

```
z_score, p_value = sm.stats.proportions_ztest([convert_old, convert_new], [n_old, n_new], alternative='smaller')
```

```
print(z_score, p_value)
```

```
1.3189241982242494 -0.9885083127980245
```

n. What are the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j, k, and l?

Put your answer here. Z-score indicates how much a given value differs from the standard deviation. The Z-score=1.31 which is less than critical value at 95% and p_value=0.985 which indicate that there is no improvement in convert for new landing page

Part III - A regression approach

ToDo 3.1

In this final part, you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

a. Since each row in the `df2` data is either a conversion or no conversion, what type of regression should you be performing in this case?

Put your answer here. This is Logistic Regression.

b. The goal is to use **statsmodels** library to fit the regression model you specified in part a above, to see if there is a significant difference in conversion based on the page-type a customer receives. However, you first need to create the following two columns in the `df2`:

```
In [134]: # intercept - It should be 1 in the entire column.
# ab_page - It's a dummy variable column, having a value 1 when an individual receives the treatment, otherwise 0.
```

```
Out[134]: df2['intercept'] = 1
df2['control', 'ab_page'] = pd.get_dummies(df2['group'])
df2.drop(['control'], axis=1, inplace=True)
df2.head()
```

```
Out[134]: user_id      timestamp      group      landing_page      converted      intercept      ab_page
0      851104    2017-01-21 22:11:48.556739    control      old_page          0          1          0
1      804228    2017-01-12 08:01:45.199739    control      old_page          0          1          0
2      804228    2017-01-12 08:01:45.199739    treatment    new_page          0          1          1
3      80584    2017-01-09 18:28:03.147656    treatment    new_page          0          1          0
4      864975    2017-01-21 01:52:26.210827    control      old_page          1          1          0
```

c. Use **statsmodels** to instantiate your regression model on the two columns you created in part (b), above, then fit the model to predict whether or not an individual converts.

```
In [135]: import statsmodels.api as sm
```

```
logit_mod = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
```

```
Out[135]: np.exp(0.8150)
```

```
Out[135]: 1.815133646015719
```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
In [135]: results = logit_mod.fit()
results.summary()
```

Optimization terminated successfully.
Current function value: 8.385118
Iterations: 6

```
Out[135]: Logit Regression Results
```

<